

DQN Walkthrough

Dan Neil

16 April 2015

Quick Paper Overview

History

Human-level control through deep reinforcement learning

Volodymyr Mnih^{1*}, Koray Kavukcuoglu^{1*}, David Silver^{1*}, Andrei A. Rusu¹, Joel Veness¹, Marc G. Bellemare¹, Alex Graves¹, Martin Riedmiller¹, Andreas K. Fiedjeland¹, Georg Ostrovski¹, Stig Petersen¹, Charles Beattie¹, Amir Sadik¹, Ioannis Antonoglou¹, Helen King¹, Dharshan Kumaran¹, Daan Wierstra¹, Shane Legg¹ & Demis Hassabis¹

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih Koray Kavukcuoglu David Silver Alex Graves Ioannis Antonoglou

Daan Wierstra Martin Riedmiller

DeepMind Technologies

{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com

Abstract



We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The model is a convolutional neural network, trained with a variant of Q-learning, whose input is raw pixels and whose output is a value function estimating future rewards. We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them.

Abstract (Following from Jakob Buhmann)



Abstract

We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The model is a convolutional neural network, trained with a variant of Q-learning, whose input is raw pixels and whose output is a value function estimating future rewards. We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them.

Why DQNs?



Show it working



Reinforcement Learning Quick Recap

- Actor chooses an action
- Action is performed and influences the environment
- Based on the new state, the system receives a reward r_t
- System should learn to choose the action that maximizes the expected discounted reward
- Challenge is to learn in spite of sparse, noisy, or temporally shifted rewards

- Q-Learning:
 - State represented by a sequence $s_t = x_1, a_1, x_2, \dots, a_t, x_t$,
 - Approximate action value function $Q(s, a)$
 - Find maximizing action $Q^*(s, a) = \max_a \mathbb{E}[R_t | s_t = s, a_t = a]$
 - Optimal solution fulfills Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$



- Approximate iteratively, s.t. $Q_i \rightarrow Q^*$, $i \rightarrow \infty$

$$Q_{i+1}(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q_i(s', a') | s, a \right]$$

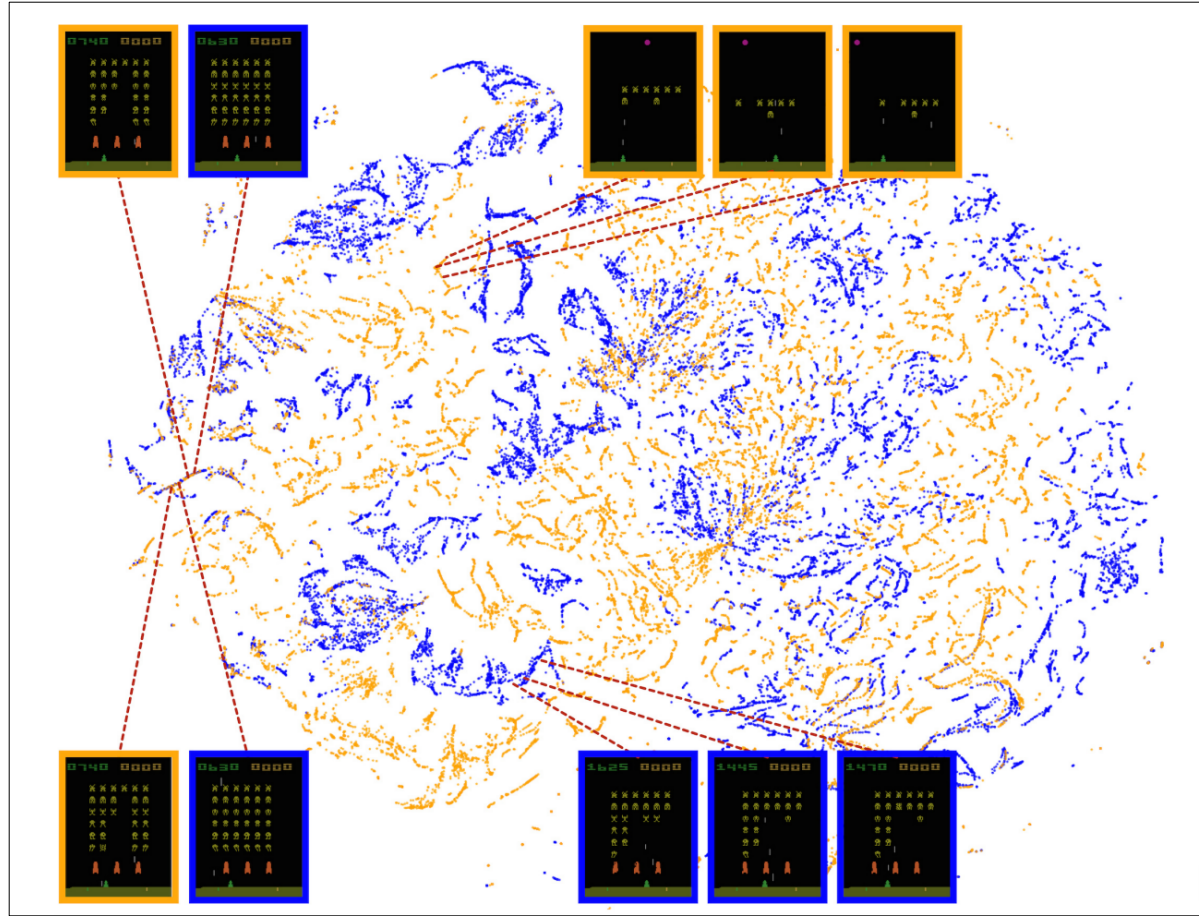
Why function approximation?

- Learn a mapping of states to true rewards
 - Sounds like neural networks (model-free generic function approximators)
- Solution: neural networks
 - Fuzzy mapping from arbitrary features to rewards
 - Many tools exist to optimize learning

Convolutional Networks



Convolutional Network Embedding



Challenges in NN as Function Approximators

- Correlations in successive states
- Small changes to Q-estimator may significantly change the policy and the subsequent data distribution
- Strong correlations exist between the action-values and the target values
- ✉ Leads to feedback loops and instability. How to solve?
- One previous approach: lots of randomly initialized networks
 - Very time costly

Enter DQNs

- Two (and a half) key ideas:
- Replay memory
 - Store some memories of the past, and randomly sample from these instead of just the current state
 - Each memory can influence training more than once (good, efficient)
 - Random sampling from past breaks consecutive correlations
 - Behaviour distribution is distributed over history
- Target is only updated periodically, allowing the action-value estimator to diverge and reduce correlations
 - No longer does an increase of $Q(s_t, a_t)$ increase $Q(s_{t+1}, a)$, avoiding feedback loop
- Clip the error rate to $(-1, 1)$

Code Deep Dive

Algorithm Overview

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

Algorithm Overview - Preprocess

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

```
170▼ function nql:preprocess(rawstate)
171▼   if self.preproc then
172       return self.preproc:forward(rawstate:float())
173       :clone():reshape(self.state_dim)
174   end
175
176   return rawstate
177 end
```

```
7   require "nn"
8   require "image"
9
10  local scale = torch.class('nn.Scale', 'nn.Module')
11
12
13▼ function scale:__init(height, width)
14    self.height = height
15    self.width = width
16 end
17
18▼ function scale:forward(x)
19    local x = x
20    if x:dim() > 3 then
21        x = x[1]
22    end
23
24    x = image.rgb2y(x)
25    x = image.scale(x, self.width, self.height, 'bilinear')
26    return x
27 end
28
29 function scale:updateOutput(input)
30     return self:forward(input)
31 end
32
33 function scale:float()
34 end
```


Algorithm Overview

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

```
315 self.transitions.add_recent_state(state, terminal)
316
317 local currentFullState = self.transitions.get_recent()
318
319 --Store transition s, a, r, s'
320 if self.lastState and not testing then
321     self.transitions.add(self.lastState, self.lastAction, reward,
322                          self.lastTerminal, priority)
323 end
324
325 if self.numSteps == self.learn_start+1 and not testing then
326     self.sample_validation_data()
327 end
328
329 curState= self.transitions.get_recent()
330 curState = curState.resize(1, unpack(self.input_dims))
331
332 -- Select action
333 local actionIndex = 1
334 if not terminal then
335     actionIndex = self:εGreedy(curState, testing_ep)
336 end
337
338 self.transitions.add_recent_action(actionIndex)
339
340 --Do some Q-learning updates
341 if self.numSteps > self.learn_start and not testing and
342    self.numSteps % self.update_freq == 0 then
343     for i = 1, self.n_replay do
344         self.qLearnMinibatch()
345     end
346 end
347
348 if not testing then
349     self.numSteps = self.numSteps + 1
350 end
351
352 self.lastState = state:clone()
353 self.lastAction = actionIndex
354 self.lastTerminal = terminal
355
356 if self.target_q and self.numSteps % self.target_q == 1 then
357     self.target_network = self.network:clone()
358 end
359
360 if not terminal then
361     return actionIndex
362 else
363     return 0
364 end
```

Algorithm Overview

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

```
315 self.transitions.add_recent_state(state, terminal)
316
317 local currentFullState = self.transitions.get_recent()
318
319 --Store transition s, a, r, s'
320 if self.lastState and not testing then
321     self.transitions.add(self.lastState, self.lastAction, reward,
322                          self.lastTerminal, priority)
323 end
324
325 if self.numSteps == self.learn_start+1 and not testing then
326     self.sample_validation_data()
327 end
328
329 curState= self.transitions.get_recent()
330 curState = curState.resize(1, unpack(self.input_dims))
331
332 -- Select action
333 local actionIndex = 1
334 if not terminal then
335     actionIndex = self:εGreedy(curState, testing_ep)
336 end
337
338 self.transitions.add_recent_action(actionIndex)
339
340 --Do some Q-learning updates
341 if self.numSteps > self.learn_start and not testing and
342    self.numSteps % self.update_freq == 0 then
343     for i = 1, self.n_replay do
344         self.qLearnMinibatch()
345     end
346 end
347
348 if not testing then
349     self.numSteps = self.numSteps + 1
350 end
351
352 self.lastState = state:clone()
353 self.lastAction = actionIndex
354 self.lastTerminal = terminal
355
356 if self.target_q and self.numSteps % self.target_q == 1 then
357     self.target_network = self.network:clone()
358 end
359
360 if not terminal then
361     return actionIndex
362 else
363     return 0
364 end
```

Algorithm Overview – Epsilon Greedy

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

```
368 function nql:eGreedy(state, testing_ep)
369     self.ep = testing_ep or (self.ep_end +
370         math.max(0, (self.ep_start - self.ep_end) * (self.ep_endt -
371             math.max(0, self.numSteps - self.learn_start))/self.ep_endt))
372     -- Epsilon greedy
373     if torch.uniform() < self.ep then
374         return torch.random(1, self.n_actions)
375     else
376         return self:greedy(state)
377     end
378 end
```

Algorithm Overview – Greedy

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

```
379
380
381 ▼ function nql:greedy(state)
382     -- Turn single state into minibatch. Needed for convolutional nets.
383 ▼     if state:dim() == 2 then
384         assert(false, 'Input must be at least 3D')
385         state = state:resize(1, state:size(1), state:size(2))
386     end
387
388     if self.gpu >= 0 then
389         state = state:cuda()
390     end
391
392     local q = self.network:forward(state):float():squeeze()
393     local maxq = q[1]
394     local besta = {1}
395
396     -- Evaluate all other actions (with random tie-breaking)
397 ▼     for a = 2, self.n_actions do
398 ▼         if q[a] > maxq then
399             besta = { a }
400             maxq = q[a]
401         elseif q[a] == maxq then
402             besta[#besta+1] = a
403         end
404     end
405     self.bestq = maxq
406
407     local r = torch.random(1, #besta)
408
409     self.lastAction = besta[r]
410
411     return besta[r]
412 end
413
```


Algorithm Overview – Get Minibatch

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

```
118 function trans:sample_one()
119     assert(self.numEntries > 1)
120     local index
121     local valid = false
122     while not valid do
123         -- start at 2 because of previous action
124         index = torch.random(2, self.numEntries-self.recentMemSize)
125         if self.t[index+self.recentMemSize-1] == 0 then
126             valid = true
127         end
128         if self.nonTermProb < 1 and self.t[index+self.recentMemSize] == 0 and
129             torch.uniform() > self.nonTermProb then
130             -- Discard non-terminal states with probability (1-nonTermProb).
131             -- Note that this is the terminal flag for s_{t+1}.
132             valid = false
133         end
134         if self.nonEventProb < 1 and self.t[index+self.recentMemSize] == 0 and
135             self.r[index+self.recentMemSize-1] == 0 and
136             torch.uniform() > self.nonTermProb then
137             -- Discard non-terminal or non-reward states with
138             -- probability (1-nonTermProb).
139             valid = false
140         end
141     end
142     return self:get(index)
143 end
144
```

Algorithm Overview – Get Rewards

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

```
-- The order of calls to forward is a bit odd in order
-- to avoid unnecessary calls (we only need 2).

-- delta = r + (1-terminal) * gamma * max_a Q(s2, a) - Q(s, a)
term = term:clone():float():mul(-1):add(1)

local target_q_net
if self.target_q then
    target_q_net = self.target_network
else
    target_q_net = self.network
end

-- Compute max_a Q(s2, a).
q2_max = target_q_net:forward(s2):float():max(2)

-- Compute q2 = (1-terminal) * gamma * max_a Q(s2, a)
q2 = q2_max:clone():mul(self.discount):cmul(term)

delta = r:clone():float()

if self.rescale_r then
    delta:div(self.r_max)
end
delta:add(q2)

-- q = Q(s, a)
local q_all = self.network:forward(s):float()
q = torch.FloatTensor(q_all:size(1))
for i=1,q_all:size(1) do
    q[i] = q_all[i][a[i]]
end
delta:add(-1, q)

if self.clip_delta then
    delta[delta:ge(self.clip_delta)] = self.clip_delta
    delta[delta:le(-self.clip_delta)] = -self.clip_delta
end

local targets = torch.zeros(self.minibatch_size, self.n_actions):float()
for i=1,math.min(self.minibatch_size,a:size(1)) do
    targets[i][a[i]] = delta[i]
end
```

Algorithm Overview – Update Target Network

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

```
356     if self.target_q and self.numSteps % self.target_q == 1 then
357         self.target_network = self.network.clone()
358     end
```


What We Can Do

DQNs for Game Playing

- For now, train a model using their default architecture
- Their parameters are identical across 49 games, indicating some robustness of these parameters
- In the future:
 - Pretrain visual model and reuse between games
 - Use model selection to pre-initialize game depending on type (puzzle, maze hack-and-slash, Mario-like, racing game, etc.)

Possible Labor Division

- Input and Preprocessing
 - Interface to the game simulator
 - Represent game state as simplified image
- DQN Network without the deep part
 - Setup epsilon-greedy framework, replay memory, and Q-target copying
 - Build generalized DQN without the q-estimator part
- ConvNet Function Approximation
 - Set up a simple interface to use the convnet to estimate the rewards

Appendix

Alg. Overview ConvNet

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ε select a random action a_t

 otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

```
9 function create_network(args)
10
11     local net = nn.Sequential()
12     net:add(nn.Reshape(unpack(args.input_dims)))
13
14     --- first convolutional layer
15     local convLayer = nn.SpatialConvolution
16
17     if args.gpu >= 0 then
18         net:add(nn.Transpose({1,2},{2,3},{3,4}))
19         convLayer = nn.SpatialConvolutionCUDA
20     end
21
22     net:add(convLayer(args.hist_len*args.ncols, args.n_units[1],
23                     args.filter_size[1], args.filter_size[1],
24                     args.filter_stride[1], args.filter_stride[1],1))
25     net:add(args.nl())
26
27     -- Add convolutional layers
28     for i=1,(#args.n_units-1) do
29         -- second convolutional layer
30         net:add(convLayer(args.n_units[i], args.n_units[i+1],
31                         args.filter_size[i+1], args.filter_size[i+1],
32                         args.filter_stride[i+1], args.filter_stride[i+1]))
33         net:add(args.nl())
34     end
35
36     local nel
37     if args.gpu >= 0 then
38         net:add(nn.Transpose({4,3},{3,2},{2,1}))
39         nel = net:cuda():forward(torch.zeros(1,unpack(args.input_dims)))
40         :cuda():nElement()
41     else
42         nel = net:forward(torch.zeros(1,unpack(args.input_dims))):nElement()
43     end
44
45     -- reshape all feature planes into a vector per example
46     net:add(nn.Reshape(nel))
47
48     -- fully connected layer
49     net:add(nn.Linear(nel, args.n_hid[1]))
50     net:add(args.nl())
51     local last_layer_size = args.n_hid[1]
52
53     for i=1,(#args.n_hid-1) do
54         -- add Linear layer
55         last_layer_size = args.n_hid[i+1]
56         net:add(nn.Linear(args.n_hid[i], last_layer_size))
57         net:add(args.nl())
58     end
59
60     -- add the last fully connected layer (to actions)
61     net:add(nn.Linear(last_layer_size, args.n_actions))
62
63     if args.gpu >= 0 then
64         net:cuda()
65     end
66     if args.verbose >= 2 then
67         print(net)
68         print('Convolutional layers flattened output size:', nel)
69     end
70     return net
71 end
72
```